



Heart disease:

From Model Building to Deployment

Let's talk about a special type of heartbreak!

● Maximilien Eyengue

Data Science Life Cycle Steps:



Business Understanding

What?

Disease

According to WHO, heart disease is the leading cause of deaths

33s

So what?

Diagnosis

Early prediction saves lives & reduces costs



Then what?

Objective

Build a tool to accurately identify patients at risk



Data Understanding

```
# download the data
```

```
!kaggle datasets download mfarhaannazirkhan/heart-dataset
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /home/maxim-eyengue/.kaggle/kaggle.json'

Dataset URL: <https://www.kaggle.com/datasets/mfarhaannazirkhan/heart-dataset>

License(s): Attribution 4.0 International (CC BY 4.0)

Downloading heart-dataset.zip to /home/maxim-eyengue/Heart-Disease-App

100%|██| 27.5k/27.5k [00:00<00:00, 216kB/s]

100%|██| 27.5k/27.5k [00:00<00:00, 215kB/s]

Five datasets with patient data combined together for heart disease diagnosis

```
# Read the dataframe
```

```
df = pd.read_csv("data/raw_merged_heart_dataset.csv")
```

```
# Three last rows
```

```
df.tail(3)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalachh	exang	oldpeak	slope	ca	thal	target
2178	59	1	3	134	204	0	1	162	0	0.8	2	2	2	0
2179	54	1	1	154	232	0	0	164	0	0.0	2	1	2	0
2180	53	1	0	110	335	0	1	143	1	3.0	1	1	3	0

Data Preparation

```
# Convert resting blood pressure
df.trestbps = pd.to_numeric(df.trestbps, errors = 'coerce')
# Convert cholesterol
df.chol = pd.to_numeric(df.chol, errors = 'coerce')
# Convert maximum heart rate
df.thalachh = pd.to_numeric(df.thalachh, errors = 'coerce')
```


No missing values? Watch out!

```
# For each categorical variable
for cat in categorical:
    # Print the variable and its values
    print(f"{cat} --> {df[cat].unique()}")

sex --> ['male' 'female']
cp --> ['asymptomatic' 'non_anginal_pain' 'atypical_angina' 'typical_angina' nan]
fbs --> ['1' '0' '?']
restecg --> ['0' '1' '2' '?']
exang --> ['0' '1' '?']
slope --> ['0' '2' '1' '?' '3']
ca --> ['0' '2' '1' '3' '4' '?']
thal --> ['1' '2' '3' '0' '?' '6' '7']
```

EDA


```
# Correlation Heatmap
plt.figure(figsize = (9,8))
plt.title("Heart Data Correlation Heatmap")
sns.heatmap(df[numerical].corr(), annot = True, linewidths = .1);
```



There is a low correlation
between numerical
features

EDA -Target Analysis

```
# Labels
labs_names = ['disease','normal']
# Target value counts
data_counts = list(df["target"].value_counts())
# Space between pies
pie_space = (0.15, 0)
# Title
plt.title('Heart Disease Repartition')
# Target variable distribution Pie chart
plt.pie(data_counts, explode = pie_space, labels = labs_names , autopct = '%1.2f%%',
        startangle = 90, shadow = True)
# Put title and plot on same axis
plt.axis('equal');
```



The dataset contains approximately the same number of healthy and sick patients

Thus, we can use accuracy score to evaluate our models

EDA – Feature Importance

```
# Apply mutual information columnwise to categorical variables
mi_scores = df_full_train[categorical].apply(mutual_info_y_score)

# Sort scores in ascending order
mi_scores.sort_values(ascending = False)
```

```
thal      0.14003
cp        0.10439
ca        0.08378
slope     0.08083
exang     0.06419
restecg   0.02542
sex       0.01781
fbs       0.00013
dtype: float64
```

The ST segment depression induced by exercise and the maximum heart rate are the most important numerical features

```
# Absolute correlations between numerical columns and target
df_full_train[numerical].corrwith(df_full_train.target).abs().sort_values(ascending = False)
```

```
oldpeak    0.362495
thalachh   0.331442
age        0.168646
chol       0.096872
trestbps   0.095350
dtype: float64
```

Among categorical features, thalassemia and chest pain information seem to be the most important to determine the disease

Data Modeling

Train – Validation – Test split:

```
# Splitting into full train and test
df_full_train, df_test = train_test_split(df, test_size = 0.2, random_state = 42)

# Splitting into train and test
df_train, df_val = train_test_split(df_full_train, test_size = 0.25, random_state = 42)

# Check datasets sizes after splitting
len(df_train), len(df_val), len(df_test)

(1083, 361, 361)
```

Target & Features:

```
# Get the target values
y_train = df_train.target.values
y_test = df_test.target.values
y_val = df_val.target.values

# Drop `target` from our data sets
del df_train["target"]
del df_test["target"]
del df_val["target"]
```

Data Modeling

Logistic Regression

81%

```
# Parameter Fine-Tuning
for param in reg_params:
    # Define the model
    model = LogisticRegression(solver = 'liblinear', C = param,
                              max_iter = 1000, random_state = 42)

    # Model fitting
    model.fit(X_train, y_train)
```

Decision Tree

97%

```
# Decision Tree fine-tuning with maximum-depth
for depth in [10, 15, 20]:
    # Decision Tree model fine-tuning with minimum samples per leaf
    for s in [1, 3, 5, 10, 15, 20, 100, 200, 500]:
        # Initialize the model with a max_depth and min_samples_leaf
        dt = DecisionTreeClassifier(max_depth = depth, min_samples_leaf = s, random_state = 42)
        # Model training
        dt.fit(X_train, y_train)
```

Random Forest

98%

```
# Fine-tuning with minimum samples per leaf
for s in min_sample:
    # Fine-tuning with number of estimators
    for n in n_params:
        # Initialize model
        rf = RandomForestClassifier(n_estimators = n,
                                   max_depth = 10,
                                   min_samples_leaf = s,
                                   random_state = 42,
                                   n_jobs = -1)

        # Model training
        rf.fit(X_train, y_train)
```

XG-Boost

97%

```
# XgBoost fine-tuning with learning rates
for eta_par in eta_values:
    # Model's parameters
    xgb_params = {
        'eta': eta_par,
        'max_depth': 10,
        'min_child_weight': 1,
        'objective': 'binary:logistic',
        'nthread': 8,
        'seed': 1,
        'verbosity': 1,
    }

    # Model training
    model = xgb.train(xgb_params, dtrain,
                     num_boost_round = 200)
```

Model Evaluation

```
# Kfold cross-validation initialization
kfold = KFold(n_splits = n_splits, shuffle = True, random_state = 1)
```

```
# For each iteration of K-fold split and the pair of indexes generated
for train_idx, val_idx in kfold.split(df_full_train):
```

```
    # Select train and validation data
    df_train = df_full_train.iloc[train_idx]
    df_val = df_full_train.iloc[val_idx]
```

```
    # Select target variables
    y_train = df_train.target.values
    y_val = df_val.target.values
```

```
    # Train model
    One_Hot_encoder, rf = train(df_train, y_train)
    # Make predictions
    y_pred = predict(df_val, One_Hot_encoder, rf)
```

```
    # Print scores' means and standard deviations
```

```
    print("Validation results:")
    print('acc mean = %.2f, acc std = +- %.2f' % (np.mean(scores), np.std(scores)))
```

```
Performing KFold Cross-Validation
```

```
Accuracy on fold 0 is 97.23 %.
```

```
Accuracy on fold 1 is 97.23 %.
```

```
Accuracy on fold 2 is 96.54 %.
```

```
Accuracy on fold 3 is 97.92 %.
```

```
Accuracy on fold 4 is 97.22 %.
```

```
Validation results:
```

```
acc mean = 97.23, acc std = +- 0.44
```

With K-Fold cross-validation, the selected model achieved an accuracy of **97.23%**, with a standard deviation of **0.44**

Model Evaluation

```
# Optimal parameters values
n_estimators, max_depth, min_samples_leaf = 40, 10, 1
# Optimal random forest model training
One_Hot_encoder, rf = train(df_full_train[categorical + numerical], df_full_train.target,
                             n_estimators = n_estimators, max_depth = max_depth,
                             min_samples_leaf = min_samples_leaf)

# Make predictions
y_pred = predict(df_test, One_Hot_encoder, rf)
# accuracy score
print('Optimal model accuracy = %.2f.' % (100 * (y_pred == y_test).mean()))

Optimal model accuracy = 95.29.
```

The accuracy of our final model is very good:

95.29%

Model Deployment



To easily create an application for our model



To encapsulate the application



For deploying the application to the cloud

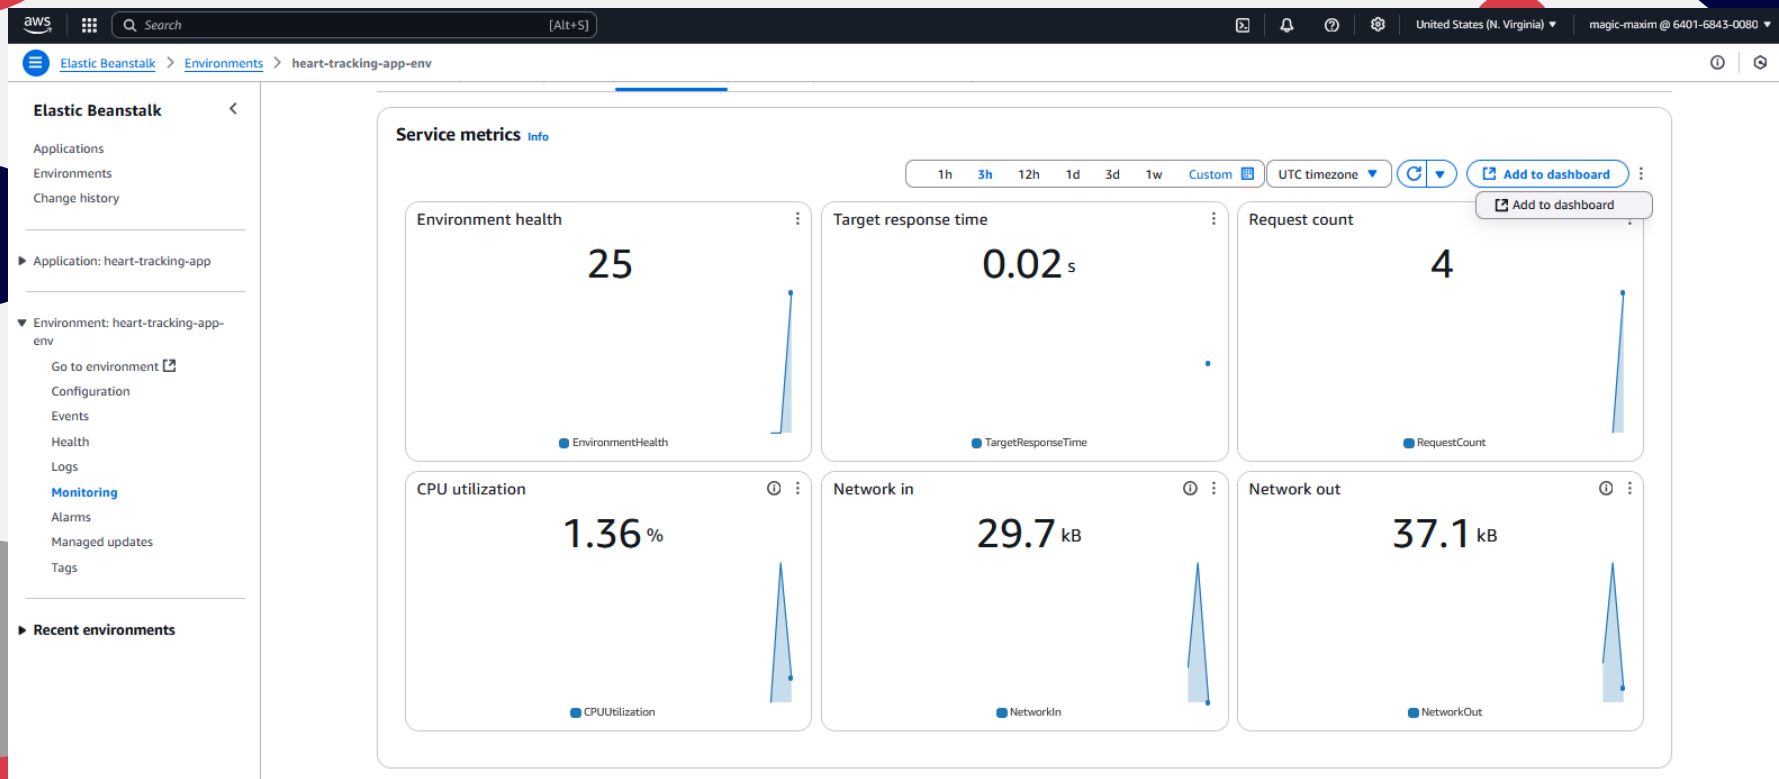
Model Deployment

```
maxim-eyengue@Magic-Maxim: ~/Heart-Disease-App
(Heart-Disease-App) (base) maxim-eyengue@Magic-Maxim:~/Heart-Disease-App$ eb create heart-tracking-app-env --enable-spot
Creating application version archive "app-ff32-250202_072749941953".
Uploading heart-tracking-app/app-ff32-250202_072749941953.zip to S3. This may take a while.
Upload Complete.
Environment details for: heart-tracking-app-env
  Application name: heart-tracking-app
  Region: us-east-1
  Deployed Version: app-ff32-250202_072749941953
  Environment ID: e-x2dyppidm
  Platform: arn:aws:elasticbeanstalk:us-east-1::platform/Docker running on 64bit Amazon Linux 2/4.0.7
  Tier: WebServer-Standard-1.0
  CNAME: UNKNOWN
  Updated: 2025-02-02 06:27:57.570000+00:00
Printing Status:
2025-02-02 06:27:55 INFO createEnvironment is starting.
2025-02-02 06:27:57 INFO Using elasticbeanstalk-us-east-1-640168430080 as Amazon S3 storage bucket for environment data.
2025-02-02 06:28:19 INFO Created security group named: sg-01777bc36e21a3314
2025-02-02 06:28:35 INFO Created target group named: arn:aws:elasticloadbalancing:us-east-1:640168430080:targetgroup/awseb-AWSEB-FQKDGZGKKWEA/82bce1c3d9e21ee0
2025-02-02 06:28:35 INFO Created security group named: awseb-e-x2dyppidm-stack-AWSEBSecurityGroup-yx2PoApVRxwh
2025-02-02 06:28:50 INFO Created Auto Scaling group named: awseb-e-x2dyppidm-stack-AWSEBAutoScalingGroup-Q3gJgruA8DYe
2025-02-02 06:28:50 INFO Waiting for EC2 instances to launch. This may take a few minutes.
2025-02-02 06:29:06 INFO Created Auto Scaling group policy named: arn:aws:autoscaling:us-east-1:640168430080:scalingPolicy:c34e8320-cfff-4b70-b994-3f9b19cce227:autoScalingGroupName/awseb-e-x2dyppidm-stack-AWSEBAutoScalingGroup-Q3gJgruA8DYe:policyName/awseb-e-x2dyppidm-stack-AWSEBAutoScalingScaleDownPolicy-xOfrrLZRHe0lw
2025-02-02 06:29:06 INFO Created Auto Scaling policy named: arn:aws:autoscaling:us-east-1:640168430080:scalingPolicy:1c8c2ed3-b156-4a97-86f1-b262b211742e:autoScalingGroupName/awseb-e-x2dyppidm-stack-AWSEBAutoScalingGroup-Q3gJgruA8DYe:policyName/awseb-e-x2dyppidm-stack-AWSEBAutoScalingScaleUpPolicy-2m1xUa603mNV
2025-02-02 06:29:06 INFO Created CloudWatch alarm named: awseb-e-x2dyppidm-stack-AWSEBCloudwatchAlarmLow-GQw3tc29x6AJ
2025-02-02 06:29:06 INFO Created CloudWatch alarm named: awseb-e-x2dyppidm-stack-AWSEBCloudwatchAlarmHigh-Zagcln0VU2LK
2025-02-02 06:30:53 INFO Created load balancer named: arn:aws:elasticloadbalancing:us-east-1:640168430080:loadbalancer/app/awseb--AWSEB-jz3DmDfnahsP/eb82d5a23d11b4b9
2025-02-02 06:30:55 INFO Created Load Balancer listener named: arn:aws:elasticloadbalancing:us-east-1:640168430080:listener/app/awseb--AWSEB-jz3DmDfnahsP/eb82d5a23d11b4b9/29aaf48dc873d25b
2025-02-02 06:31:47 INFO Instance deployment completed successfully.
2025-02-02 06:32:00 INFO Application available at heart-tracking-app-env.eba-qzigkeuz.us-east-1.elasticbeanstalk.com.
2025-02-02 06:32:01 INFO Successfully launched environment: heart-tracking-app-env

(Heart-Disease-App) (base) maxim-eyengue@Magic-Maxim:~/Heart-Disease-App$

maxim-eyengue@Magic-Maxim: ~/Heart-Disease-App
(ml-zoomcamp) maxim-eyengue@Magic-Maxim:~/Heart-Disease-App$ python predict_test_cloud.py
{'heart_disease': False, 'heart_disease_probability': 0.3}
The patient seems healthy: no treatment needed.
(ml-zoomcamp) maxim-eyengue@Magic-Maxim:~/Heart-Disease-App$
```

Model Deployment





Key Takeaways



What to do

- Be curious, judgmental and argumentative
- Think on how to improve



What not to do

- Rush into projects
- Skip data cleaning

There is always room for improvement...



?

Thanks!

Do you have any questions?
maximilien.boulou@aims-cameroon.org
+237 699 05 45 37



CREDITS: This presentation template was created by
Slidesgo, including icons by **Flaticon**, and
infographics & images by **Freepik**

Thanks also to the DataTalks.Club
[Github Project Link](#)